

# All Relevant Feature Selection

---



# Racapping Our Path So Far

## **We started with a somewhat clear goal**

- Given data containing candidate correlates and a discrete target
- ...We aimed at identifying the most relevant correlates

## **We applied a baseline approach (Lasso) to:**

- Obtain a surrogate for our data-generation process
- Analyze the impact of each candidate correlate (feature)
- Identify the most relevant correlates

## **Our baseline turned out to be largely insufficient, so we:**

- Trained a non-linear model to obtain a more reliable surrogate
- Learn to assess importance via a permutation-based method
- Learned to explain individual examples via SHAP

 \*We still have a couple of major open problems...

# Open Problems

## **There's a mismatch between local and global explanations**

- We are using SHAP to assess local feature effects
- ...And permutation importance for global feature effects

As a side effect, there may be inconsistencies in our analysis

## **We still don't know how to identify the most relevant features**

- Like in the Lasso approach we could think of using a threshold
- ...But we still don't know how such threshold should be calibrated

**It's time that we fix both of them**



# Global Feature Analysis via SHAP

## SHAP explanations can be aggregated to get global importance scores

By default, this is done by averaging absolute SHAP values:

$$\bar{\phi}_j(x) = \frac{1}{n} \sum_{i=1}^m |\phi_j(x_i)|$$

- Other aggregation functions can also be used (e.g. max)

## By using aggregated SHAP scores

...We ensure that our local and global analysis have a similar semantic

- Permutation Importance are a perfectly viable approach
- ...And sometimes may be more appropriate than SHAP

However, when doing a rigorous analysis consistency is important

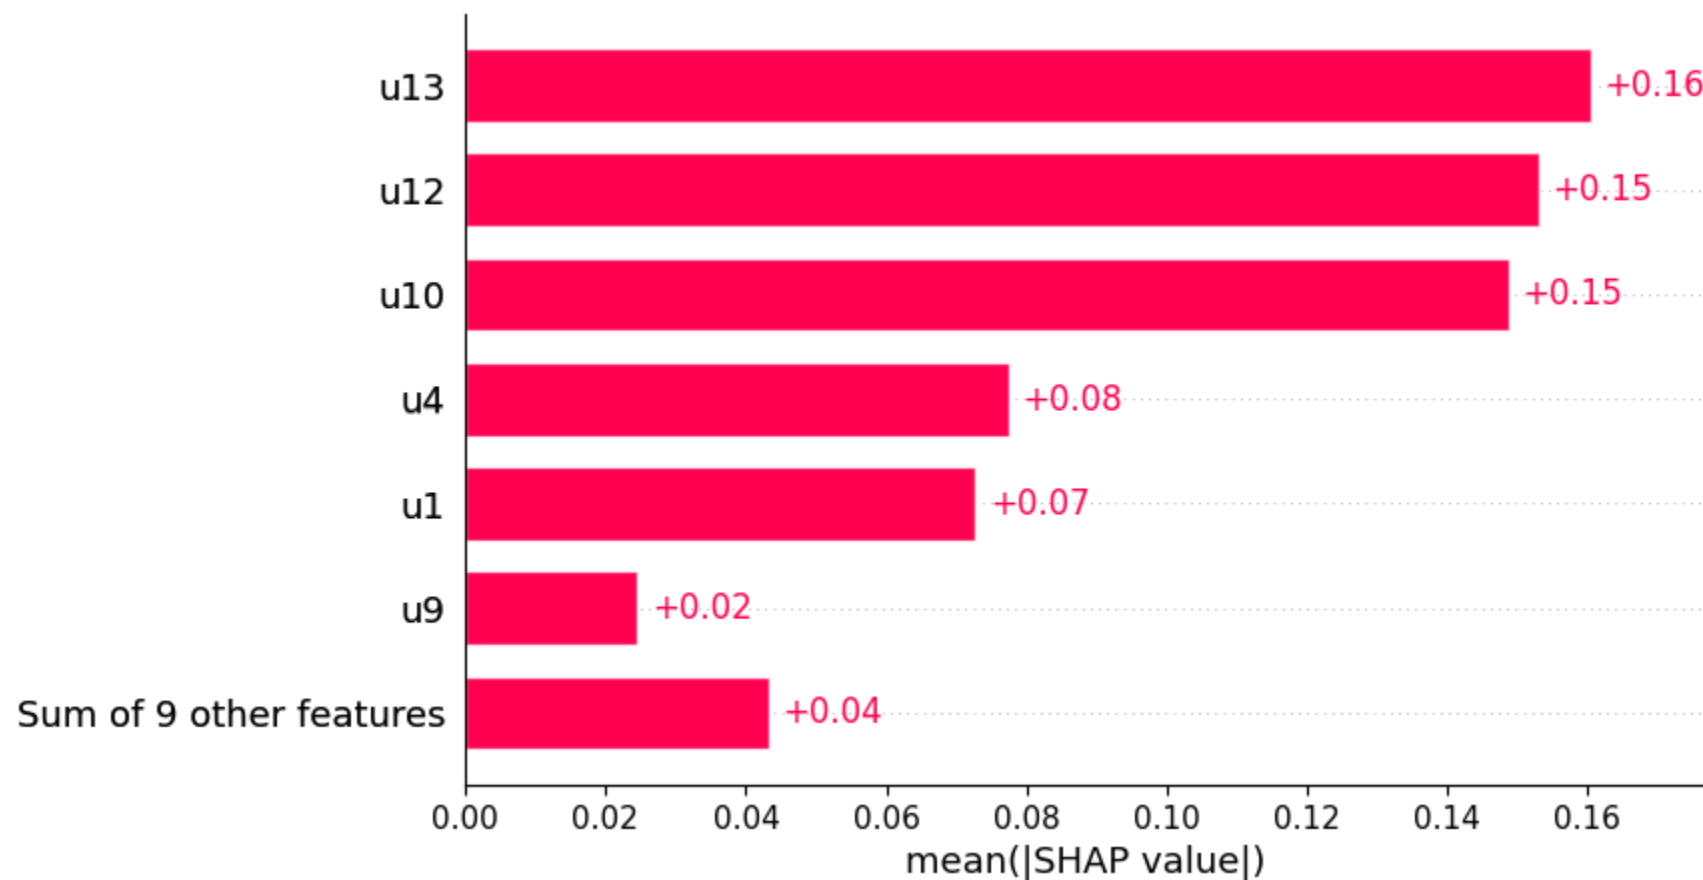


# Global Feature Analysis via SHAP

The SHAP library provide convenience functions to plot aggregated values

Here's how to plot mean (absolute) SHAP values:

```
In [9]: shap.plots.bar(shap_values, max_display=7)
```

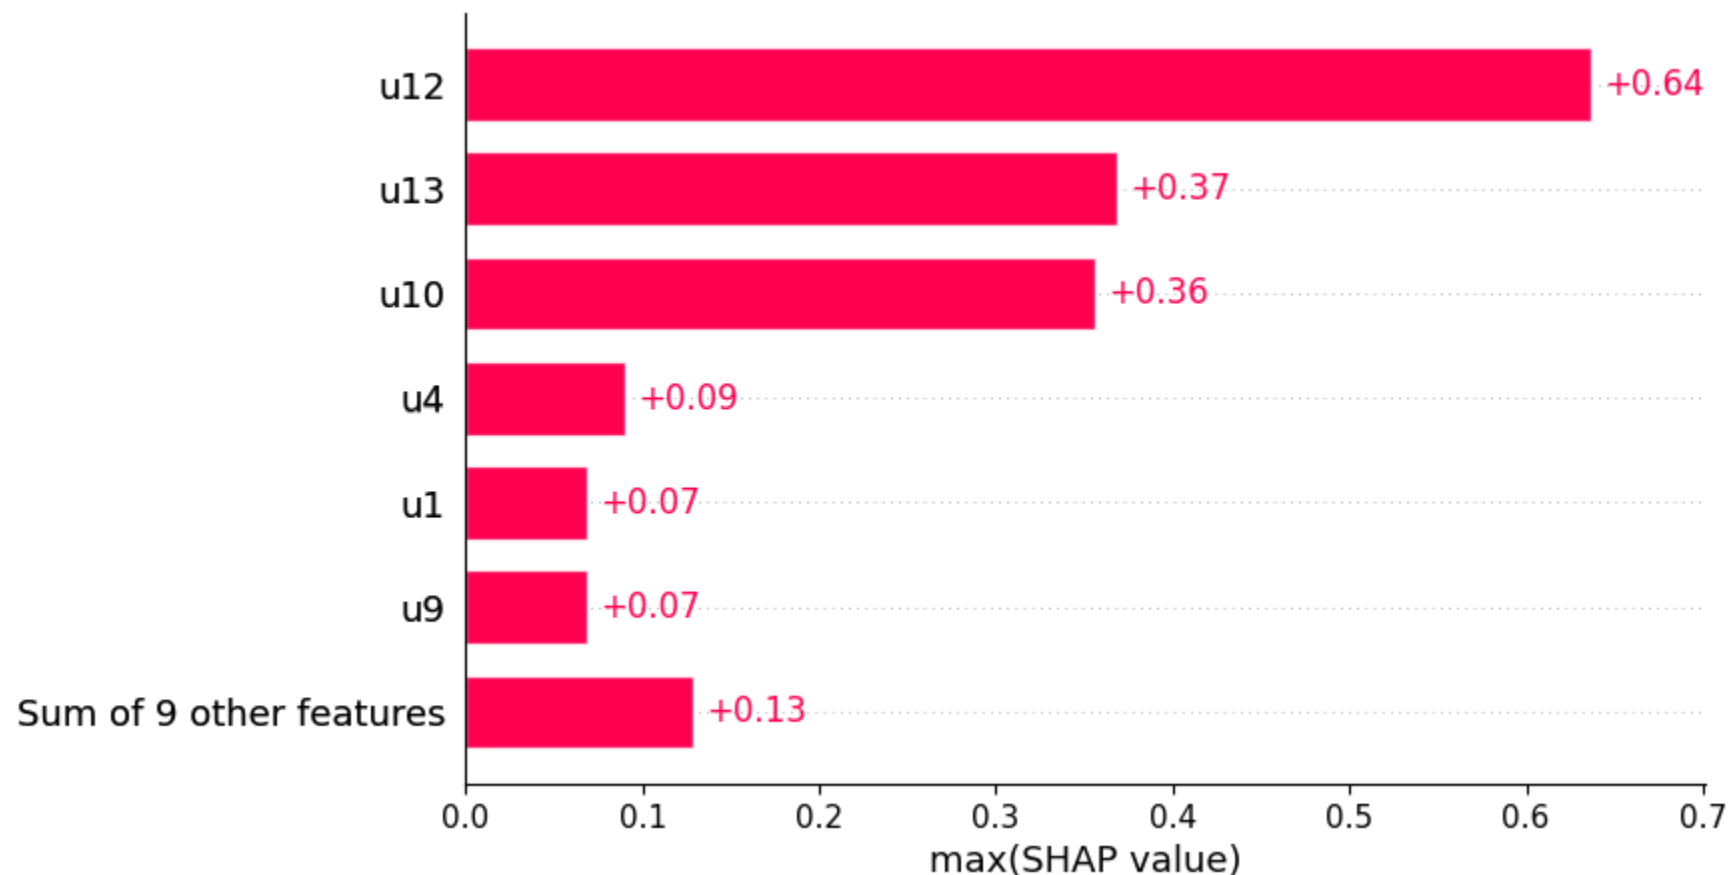


# Global Feature Analysis via SHAP

The SHAP library provide convenience functions to aggregated values

Here's how to display the maximum (absolute) SHAP values:

```
In [10]: shap.plots.bar(shap_values.max(0), max_display=7)
```



# Semantics for Feature Selection

A viable approach for feature selection consists in solving:

$$\operatorname{argmin}_{S \subseteq \mathcal{X}} \left\{ |S| : \hat{y} = \hat{f}_S(x_S), L(y, \hat{y}) \leq \theta \right\}$$

Where  $\mathbf{x}$ ,  $\mathbf{y}$  denote all the training data. Intuitively:

- We search for the smallest subset of features  $\mathcal{S}$
- ...Such that a model  $\hat{f}_S$  trained over only over them
- ...Still has an acceptable (cross-validation) accuracy

Heuristics (e.g. greedy search) can be used to improve scalability



# Semantics for Feature Selection

A viable approach for feature selection consists in solving:

$$\operatorname{argmin}_{S \subseteq \mathcal{X}} \left\{ |S| : \hat{y} = \hat{f}_S(x_S), L(y, \hat{y}) \leq \theta \right\}$$

Where  $\mathbf{x}$ ,  $\mathbf{y}$  denote all the training data. Intuitively:

- We search for the smallest subset of features  $\mathcal{S}$
- ...Such that a model  $\hat{f}_S$  trained over only over them
- ...Still has an acceptable (cross-validation) accuracy

Heuristics (e.g. greedy search) can be used to improve scalability

## This optimization-driven approach

- ...Can be customized by adjusting the constraint and cost function



- ...Can reduce data storage and location costs on the deployed model



# Semantics for Feature Selection

If we care just about **cost and accuracy**, the optimization approach is perfect

But it is not suitable for our current case study...  
Can you tell why?



# Semantics for Feature Selection

If we care just about **cost and accuracy**, the optimization approach is perfect

But it is not suitable for our current case study...  
Can you tell why?

**For a number of reasons:**

- We care about finding **all the relevant features**, not a minimal set
- How should the accuracy threshold be calibrated?
- What about the noise induced by retraining?



# Semantics for Feature Selection

If we care just about **cost and accuracy**, the optimization approach is perfect

But it is not suitable for our current case study...  
Can you tell why?

**For a number of reasons:**

- We care about finding **all the relevant features**, not a minimal set
- How should the accuracy threshold be calibrated?
- What about the noise induced by retraining?

If we wish to use ML for data analysis, we need another approach

**...In particular, we will rely on statistical hypothesis testing (HT)**



# Statistical Hypothesis Testing

**HT is one of the main mathematical workhorses of scientific research**

HT builds evidence for a hypothesis by **refuting** a competing one:

- We start from a hypothesis  $H$  and some data  $x$
- We formulate a competing **null hypothesis**  $H_0$
- We define a **test statistic**  $T(X)$ , monotonically related to  $H$
- We define the theoretical probability of  $T(X)$  under  $H_0$ , i.e.  $P(T(X) \mid H_0)$
- We compute the its empirical value for our data  $t = P(T(x) \mid H_0)$
- We compute the probability that  $T(X)$  is as extreme as  $t$  under  $H_0$ , i.e.:

$$p = P(T(X) \geq t \mid H_0)$$

- If  $p \leq 1 - \alpha$  for some confidence  $\alpha$ , we **reject the null hypothesis**



This is probably very confusing...

**Let's make an example for our case**



# Hypothesis, Data, and Null-Hypothesis

First, we need to define our **hypothesis** and **data**

We care about identifying correlates, so a good choice might be:

- $H \equiv "r(X, Y) \geq r^*", \text{ for some correlation measure } r$
- $\text{data} \equiv "x, y", \text{ i.e. our sample}$



# Hypothesis, Data, and Null-Hypothesis

First, we need to define our **hypothesis** and **data**

We care about identifying correlates, so a good choice might be:

- $H \equiv "r(X, Y) \geq r^*",$  for some correlation measure  $r$
- $\text{data} \equiv "x, y",$  i.e. our sample

Now we need a competing **null hypothesis**

A good choice might be  $H_0 \equiv \text{"the observed result is due to chance"}$

- In most cases, the null hypothesis assumes what we observe is due to chance
- If we manage to reject it, we can claim that  $H$  is more likely true
- The tricky part is choosing a  $H_0$  for which we can compute probabilities
- ...Without introducing **unnecessary assumptions**



Now we need some "test" related to  $H$  and  $H_0$   
..And it must be something for which we can compute probabilities

**How do we do that?**





# Test Statistic and Theoretical Probability

Let's consider the event " $r(X, Y) \geq r^*$ "

Since it has a binary outcome, it will follow a Bernoulli distribution

- If we assume that the correlation is due to chance...
- ...Then the associated probability should be  $1/2$

Let's pretend we make repeated experiments

The number of observe events  $T(X, Y)$  will follow a binomial distribution

- Given the number of experiments  $n$
- ...The probability of  $T(X, Y)$  should be  $P(T \mid H_0) = B(n, 1/2)$

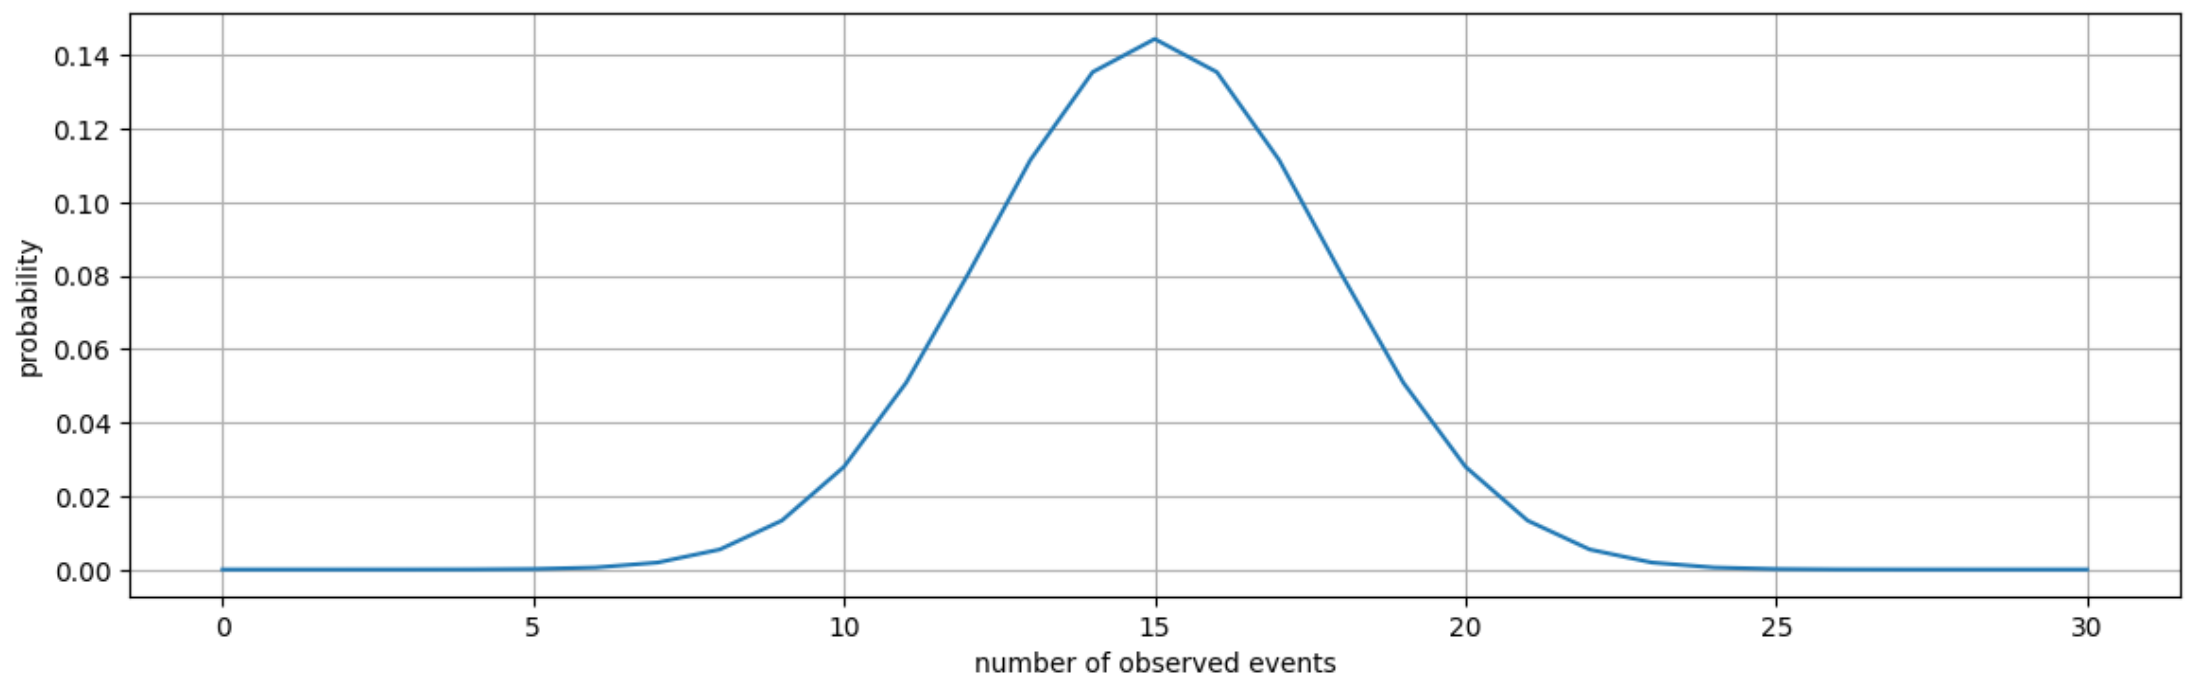
$T$  is our **test statistics**,  $B(n, 1/2)$  its theoretical distribution under  $H_0$



# Theoretical Probability Computation

We can easily compute and plot the distribution

```
In [41]: util.binomial_plot(n=30, p=0.5, figsize=figsize)
```



- This tells us how likely we are to observe a certain number of events
- ...Assuming that the null hypothesis is true



# What about the empirical probability?

We need to simulate lack of correlation  
...Without additional assumptions



# Empirical Probability Computation

## We can use a Monte-Carlo approach

The trick is once again relying on **permutations**

- If we shuffle the values of one variable (say the values  $\mathbf{x}$  of  $\mathbf{X}$ )
- ...We can get a correlation with  $\mathbf{Y}$  only by chance
- ...But we otherwise preserve the distribution of the sample

## We can mitigate sampling noise via repeated experiments

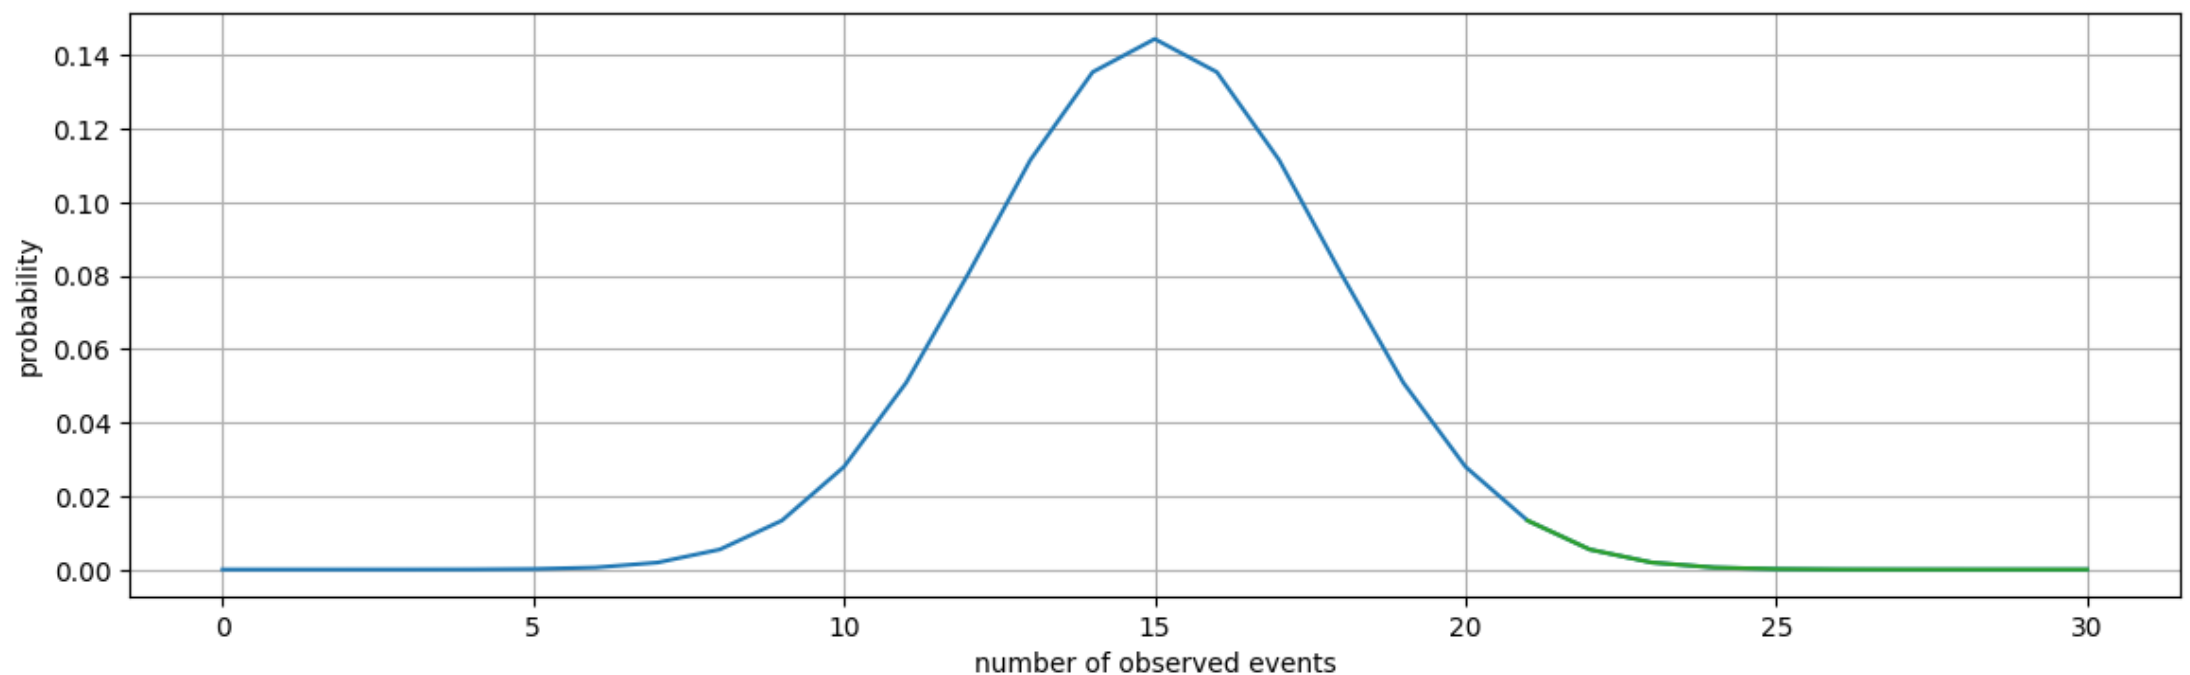
- Then, we take our **empirically observed** number of events  $t = T(\mathbf{x}, \mathbf{y})$
- ...And we match it against the theoretical probability
- We care about the probability that  $T(\mathbf{X}, \mathbf{Y}) \geq t$
- ...Since any value larger than  $t$  would still support the null-hypothesis



# $p$ -Value and the Statistical Test

Basically, there is a "target interval" in the distribution

```
In [42]: util.binomial_plot(n=30, p=0.5, r_alpha=0.01, figsize=figsize)
```



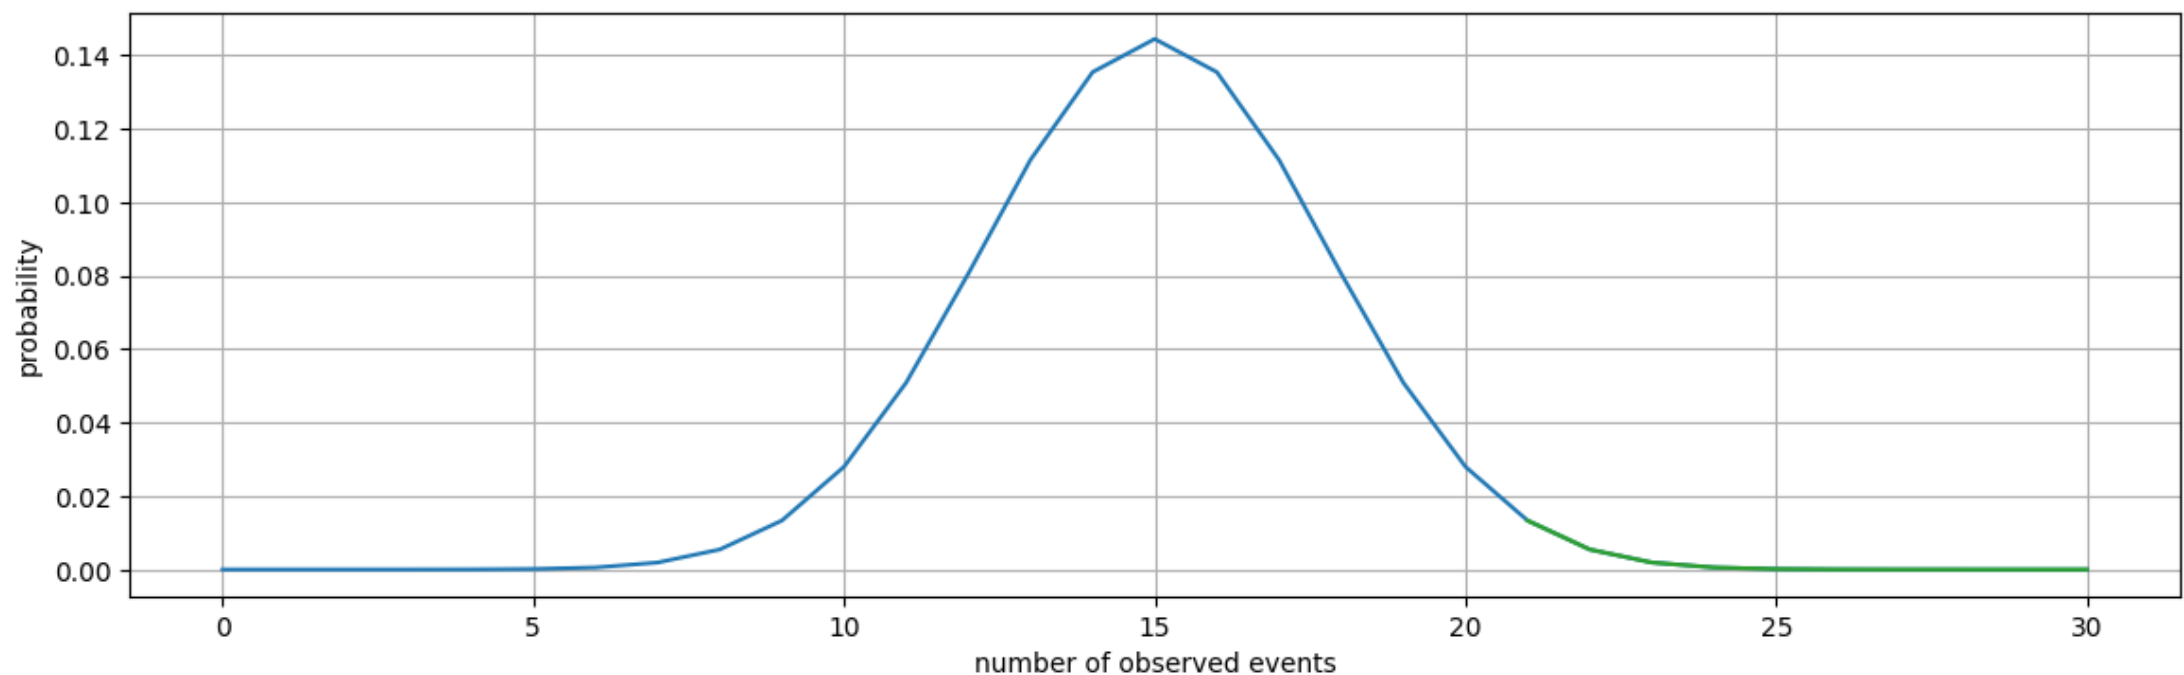
- For any value in the interval, we have  $P(T(X, Y) \geq t \mid H_0) \leq 1 - \alpha$
- ...Where  $\alpha$  is our desired confidence level



# $p$ -Value and the Statistical Test

Basically, there is a "target interval" in the distribution

```
In [44]: util.binomial_plot(n=30, p=0.5, r_alpha=0.01, figsize=figsize)
```



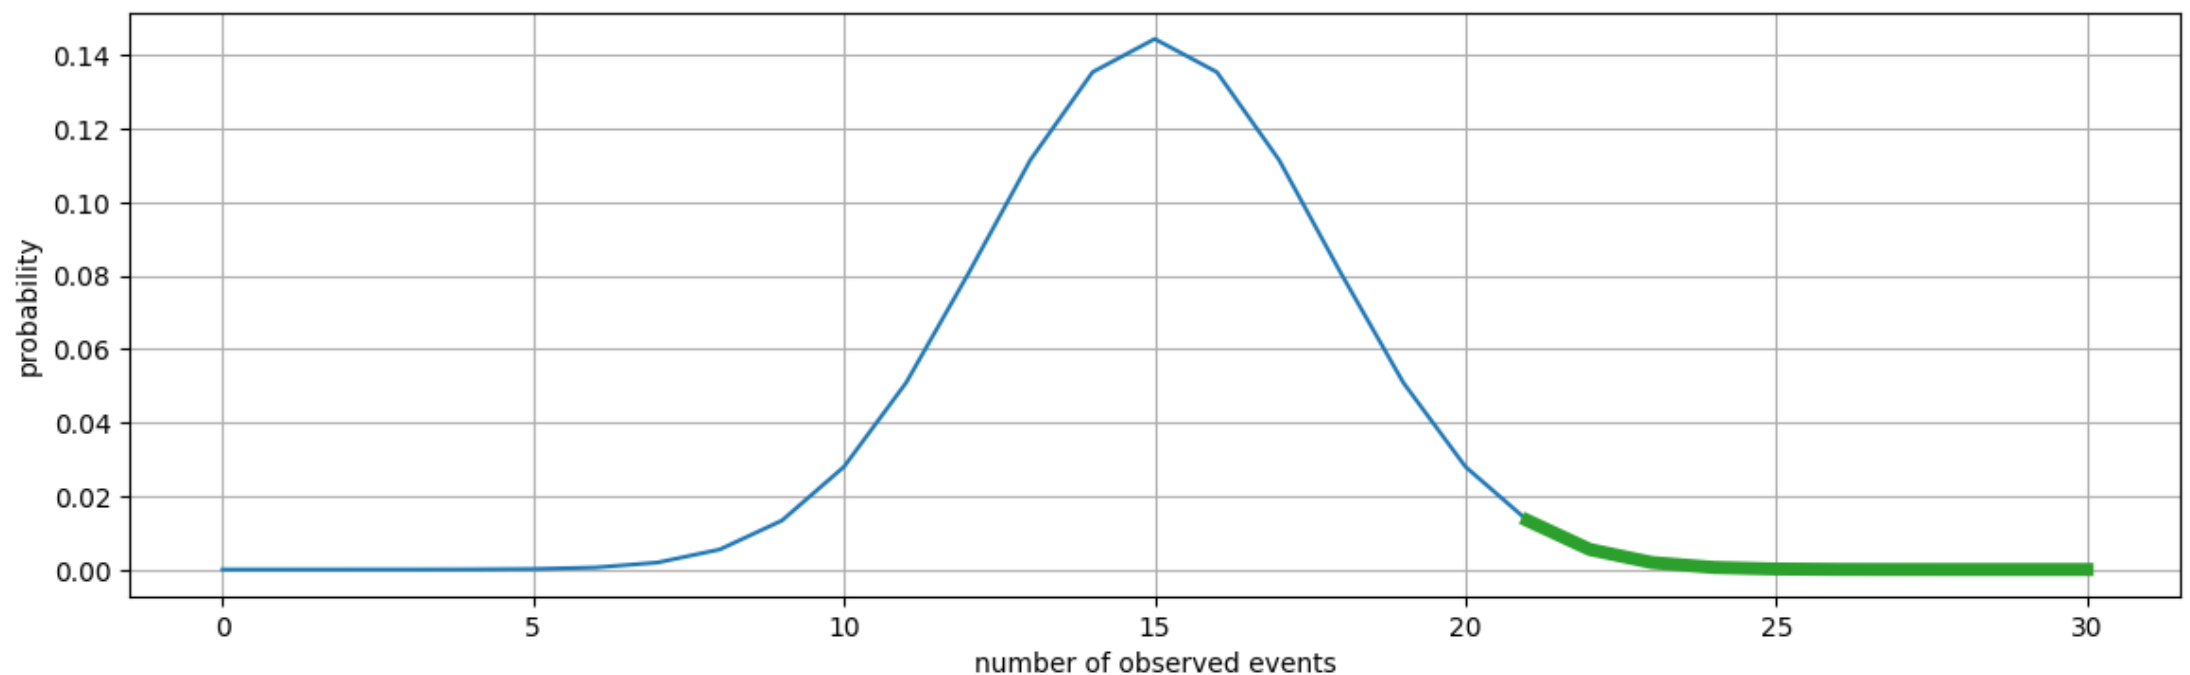
- We still need a threshold (i.e.  $\alpha$ ) to define the interval
- ...But it's a probability, so it's easier to define (usually  $\alpha = 0.01$  or  $\alpha = 0.05$ )



# $p$ -Value and the Statistical Test

Basically, there is a "target interval" in the distribution

```
In [50]: util.binomial_plot(n=30, p=0.5, r_alpha=0.01, figsize=figsize)
```



- In practice it's more common to compute the  $p$ -value  $P(T(X, Y) \geq t \mid H_0)$
- ...Which can then be immediately compared with  $1 - \alpha$



# Back to the Procedure Description

The procedure should be clearer now

Let's recap the steps:

- We start from a hypothesis  $H$  and some data  $x$
- We formulate a competing null hypothesis  $H_0$
- We define a test statistic  $T(X)$ , monotonically related to  $H$
- We define the theoretical probability of  $T(X)$  under  $H_0$ , i.e.  $P(T(X) \mid H_0)$
- We compute the its empirical value for our data  $t = P(T(x) \mid H_0)$
- We compute the probability that  $T(X)$  is as extreme as  $t$  under  $H_0$ , i.e.:

$$p = P(T(X) \geq t \mid H_0)$$

- If  $p \leq 1 - \alpha$  for some confidence  $\alpha$ , we reject the null hypothesis





# Testing a Hypothesis and Its Negation

In our case, the method works also for testing **lack of correlation**

- Our hypothesis becomes  $\neg H \equiv "r(X, Y) < r^*"$
- The null hypothesis is the same as before
- The test statistics is just  $-T(X, Y)$ , for the same  $T$  as before

Then we can proceed as in the previous case

**Since we are relying on the same test statistics**

...We can use the same set of experiments to test both hypotheses

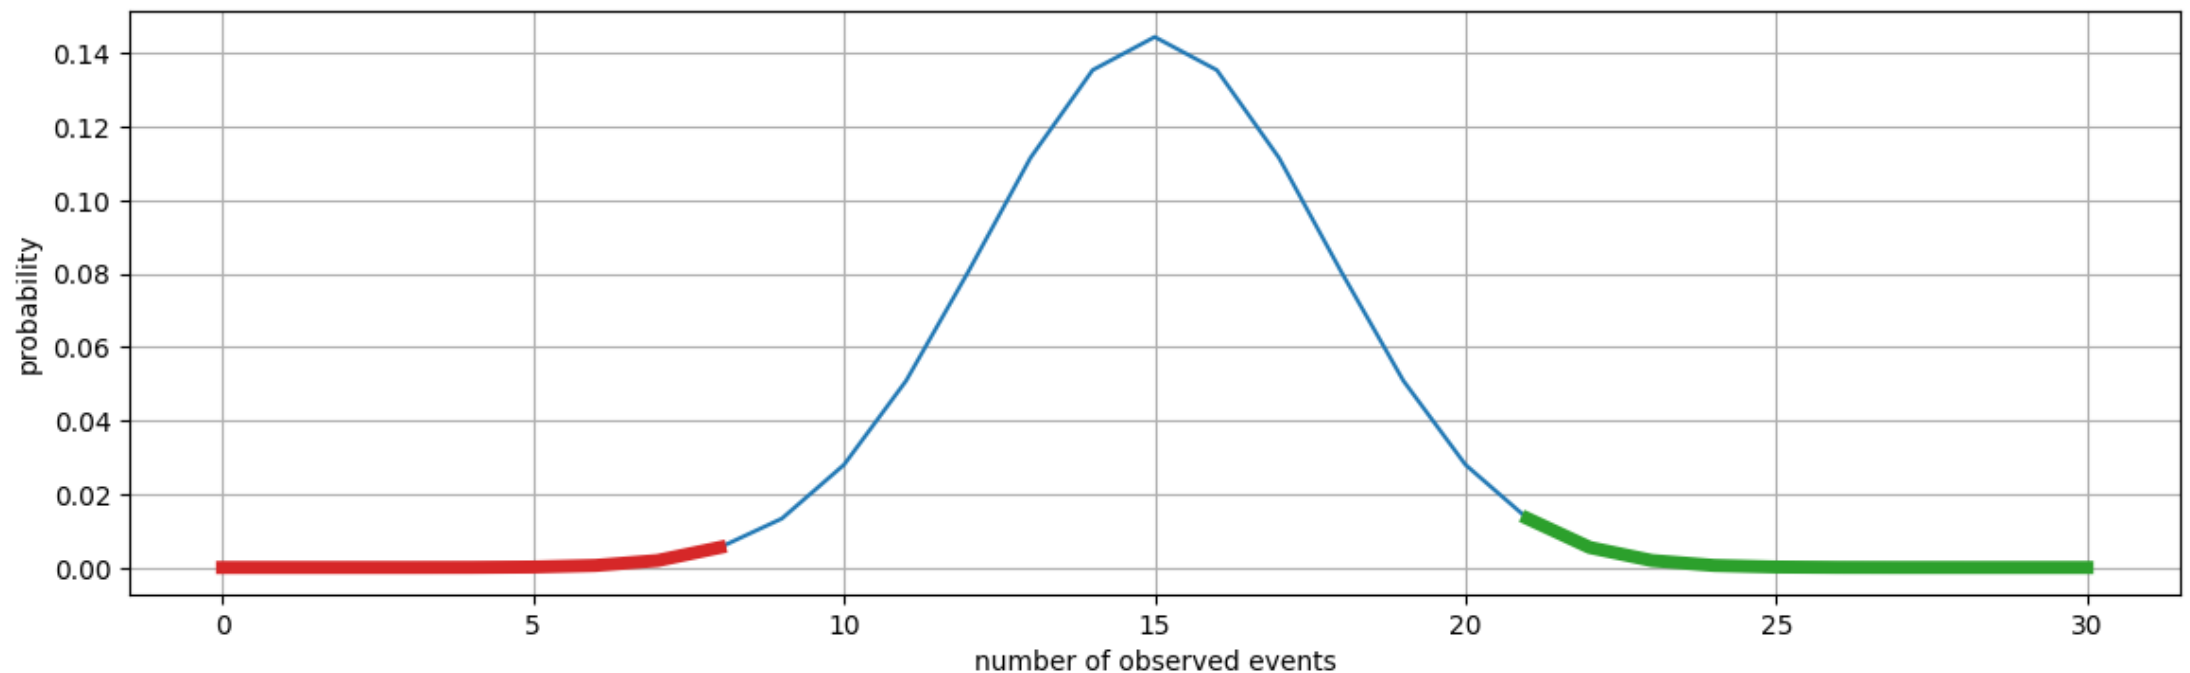
- Intuitively, in both cases we look at the number of events  $T(xy)$
- If  $T(x, y)$  is sufficiently high, it's likely that the  $H$  holds
- ...If  $T(x, y)$  is sufficiently low, it's likely that  $\neg H$  holds



# Testing a Hypothesis and Its Negation

In other words, we will end up having **two target intervals**

```
In [51]: util.binomial_plot(n=30, p=0.5, l_alpha=0.01, r_alpha=0.01, l_color='tab:red', r_color='tab:')
```



- If  $T(x, y)$  lands in the green region, we support  $H$  (e.g. correlation)
- If  $T(x, y)$  lands in the red region, we support  $\neg H$  (e.g. lack of correlation)
- If  $T(x, y)$  lands in the center region, we support no claim

# Boruta

**The approach we have just seen is the backbone of the Boruta algorithm**

- The Boruta algorithm is a SotA feature selection method
- ...That relies in statistical HT to determine relevant features

**Like in our analysis, the method relies on surrogate models**

...And in particular on tree ensembles (the name refers to a Slavic forest spirit)

- As a consequence, the algorithm can deal with non-linear correlations
- ...And accounts for interactions between multiple features

**Boruta is an all-relevant feature selector**

- This makes it particularly well suited for scientific analyses
- ...But it can be used to reduce data collection costs or improve generalization



# Statistical Testing in Boruta

## Boruta relies on a measure of feature importance

- The original algorithm and the BorutaPy\_package use permutation importance
- The more recent BorutaShap\_package relies on average SHAP values

Let  $\phi_j(x, y)$  be the importance of feature  $j$ , on a reference dataset  $(x, y)$

## The hypothesis being tested consists of:

$$\phi_j((x, \tilde{x}), y) > \max_{j \in \tilde{\mathcal{X}}} \phi_j((x, \tilde{x}), y)$$

- The dataset is **augmented** by introducing permuted versions of all features
- These are referred to as **shadow features**
- $\tilde{X}$  refers to the set of such shadow features and  $\tilde{x}$  to their values



# Statistical Testing in Boruta

The testing statistics  $T$  is similar to the one we used:

- The algorithm performs multiple experiments (retraining the model)
- ...And counts the number of times the hypothesis is satisfied (or "hits")

The theoretical distribution is **mostly** a binomial

- The algorithm needs to apply some statistical corrections
- ...Since we are testing multiple features together (we have a **max**)

**Boruta tests both the positive and negative hypothesis**

Therefore, at the end of the process:

- Some features will be **confirmed important**
- Some features will be **confirmed unimportant**

-  ■ Some features will remain **tentative**

# Using Boruta in Practice

## We'll use Boruta through the BorutaShap package

```
In [59]: bfs = BorutaShap(importance_measure='shap', classification=True)
bfs.fit(X=X, y=y, n_trials=100, sample=False, train_or_test='test', normalize=True, verbose=

100% 100/100 [01:15<00:00, 1.28it/s]

4 attributes confirmed important: ['u1', 'u12', 'u10', 'u13']
11 attributes confirmed unimportant: ['u9', 'u6', 'u3', 'u0', 'u2', 'u14', 'u11', 'u7', 'u
4', 'u5', 'u8']
0 tentative attributes remains: []
```

- We can choose to use either the testing or training importance
- The algorithm also determines the best number of estimators
- The algorithm allow the use of a clever sampling procedure
- ...To reduce the number of averaged SHAP values (and therefore the run-time)

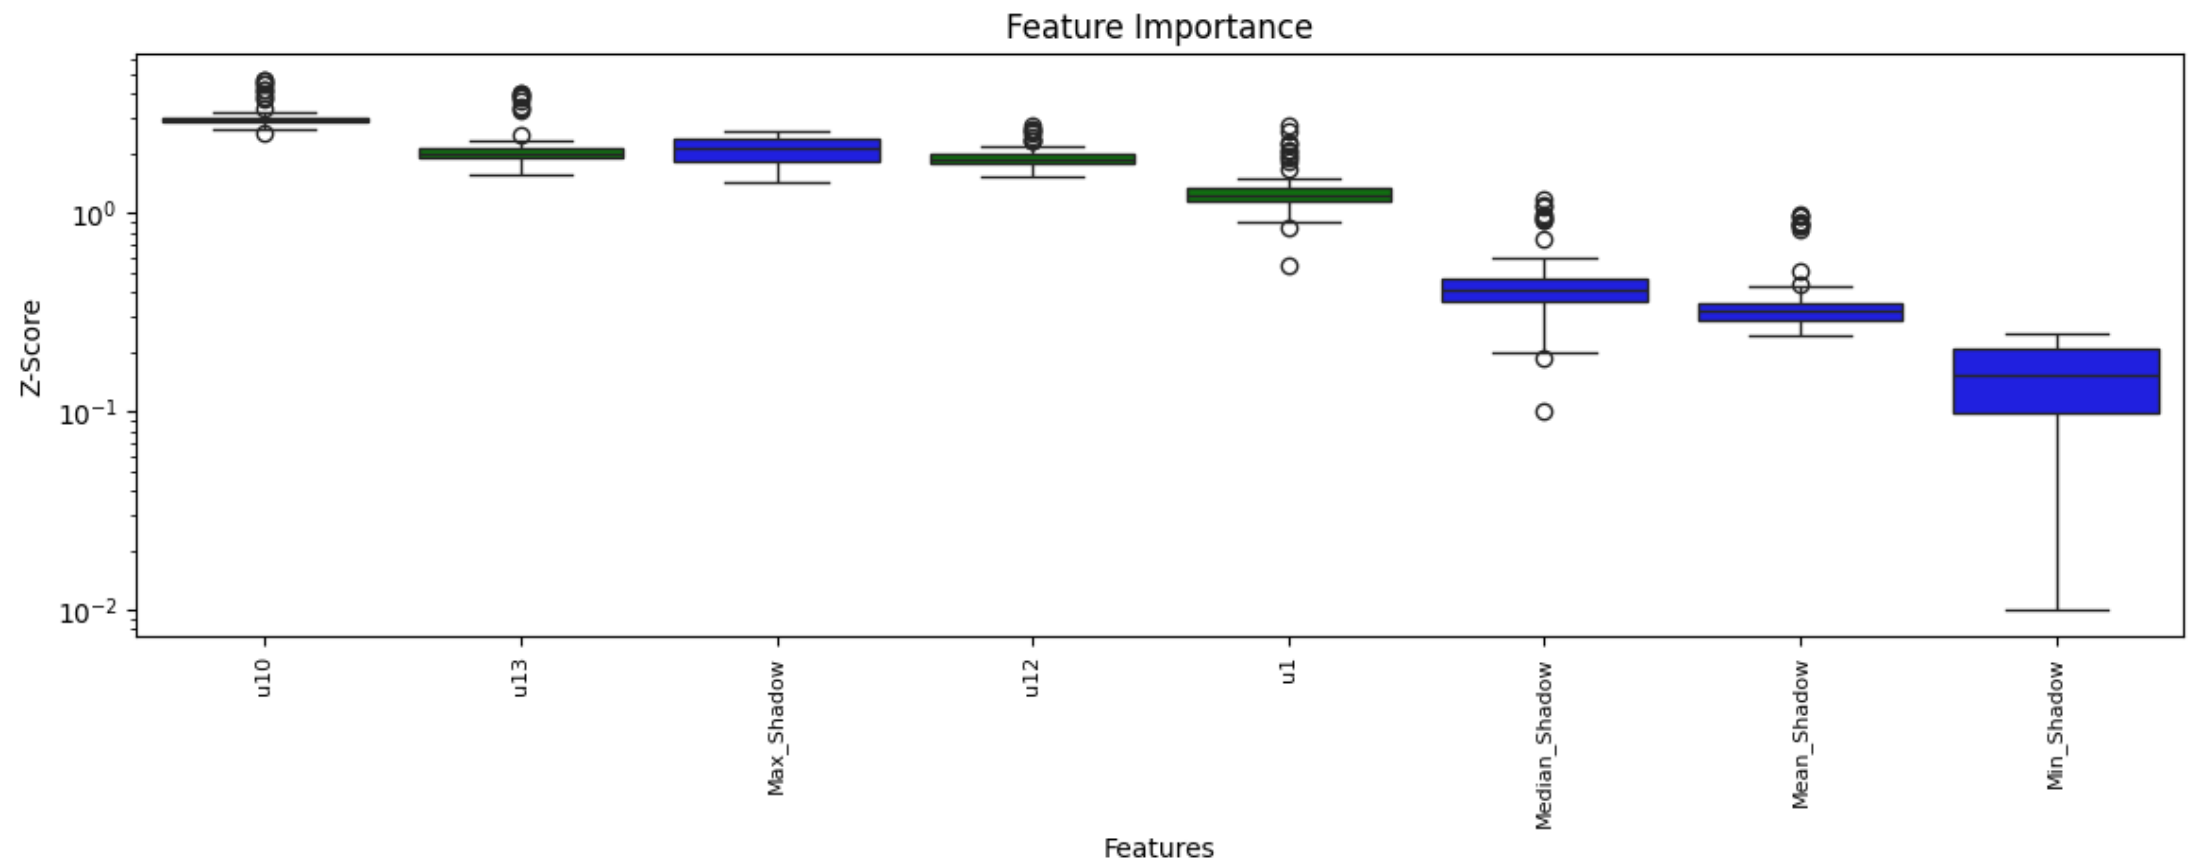
**Warning:** as of Nov 2023, the PyPI version of BorutaShap is not compatible with the most recent scikit-learn release (this lecture uses the GitHub version)



# Using Boruta in Practice

We can plot the  $\phi_j$  distribution for the **confirmed important** features

```
In [58]: bfs.plot(which_features='accepted', figsize=figsize)
```

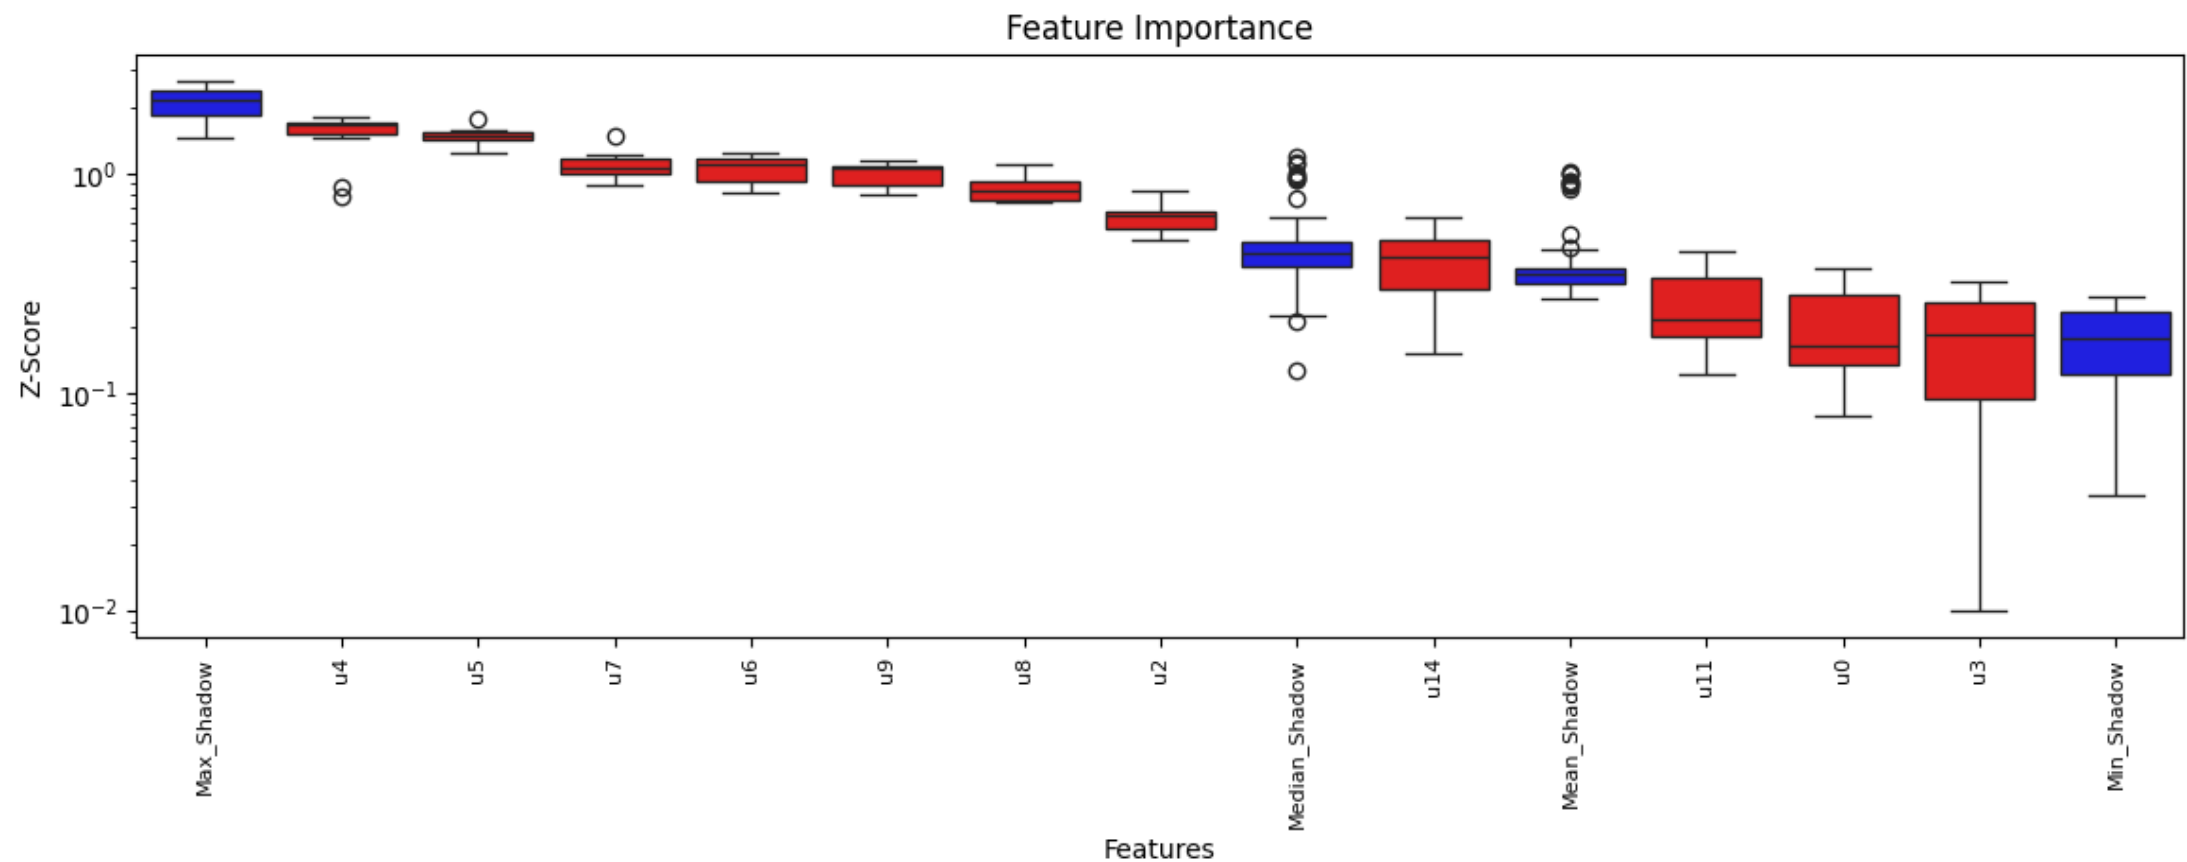


Distribution data for the shadow features is shown for comparison

# Using Boruta in Practice

We can do the same for the **confirmed unimportant** features

```
In [60]: bfs.plot(which_features='rejected', figsize=figsize)
```



 Rejected features are shown in red



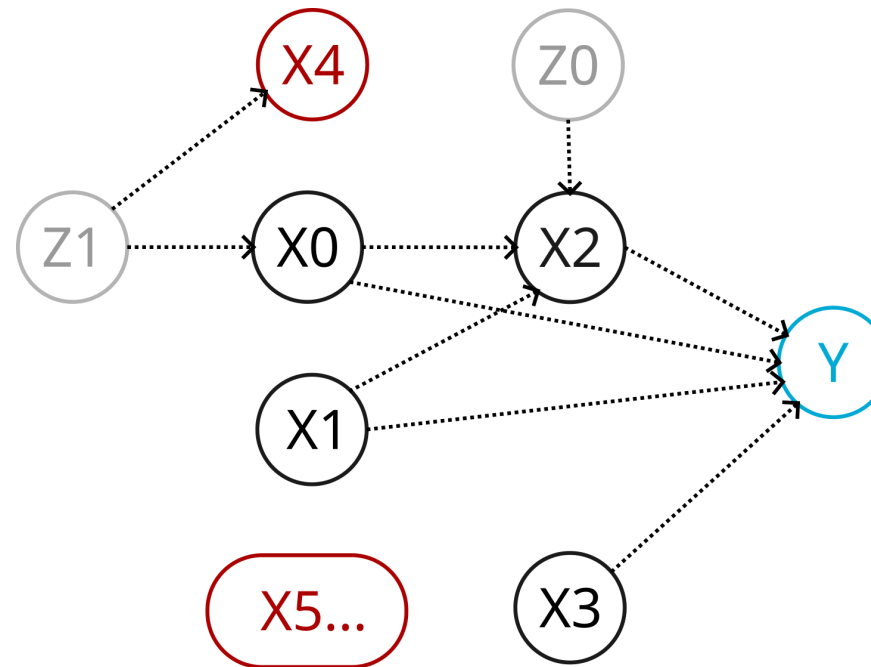
# Ok, but... Did it work?

In our controlled setting, we can inspect the ground truth process



# Checking the Ground Truth

The ground-truth process is described by this causal graph:

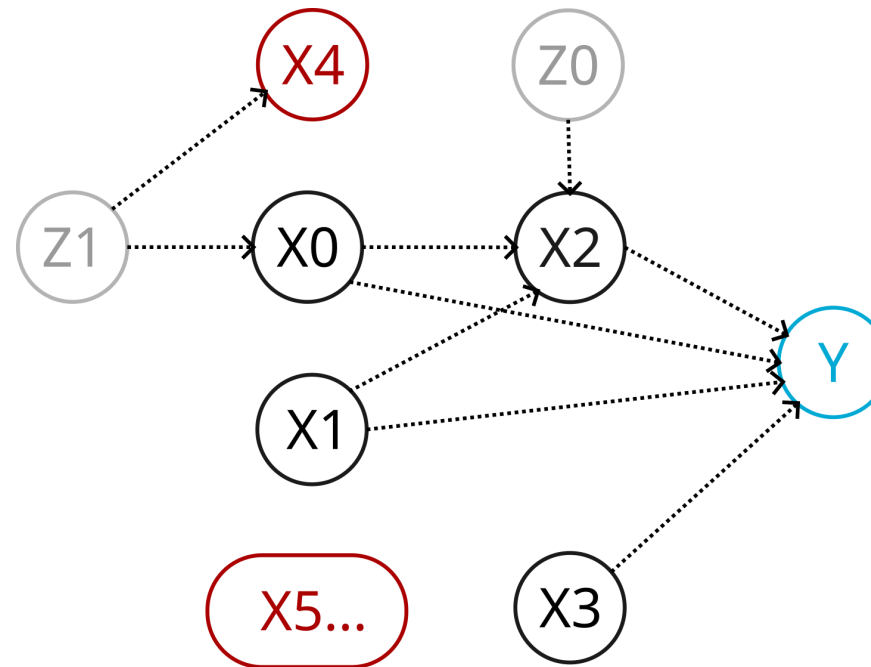


- The  $Y$  variable (in **blue**) is the target
- The variables in **black** are those that are relevant
- The variables in **gray** are not observable, i.e. latent
- The variables in **red** are irrelevant



# Checking the Ground Truth

The process was engineered to contain several classical cases



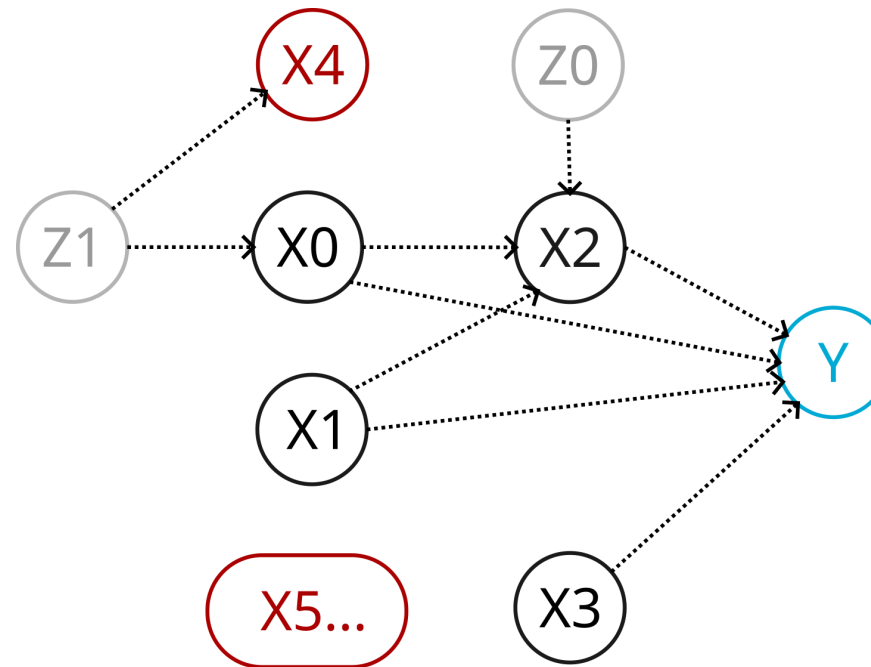
$X_2$  is a **mediator** between  $X_0$ ,  $X_1$  and  $Y$

- The variable partially hides the effect of  $X_0$  and  $X_1$
- If it does that completely, even Boruta cannot mark  $X_0$  and  $X_1$  as important
- Depending on the use case, this might be an issue



# Checking the Ground Truth

The process was engineered to contain several classical cases



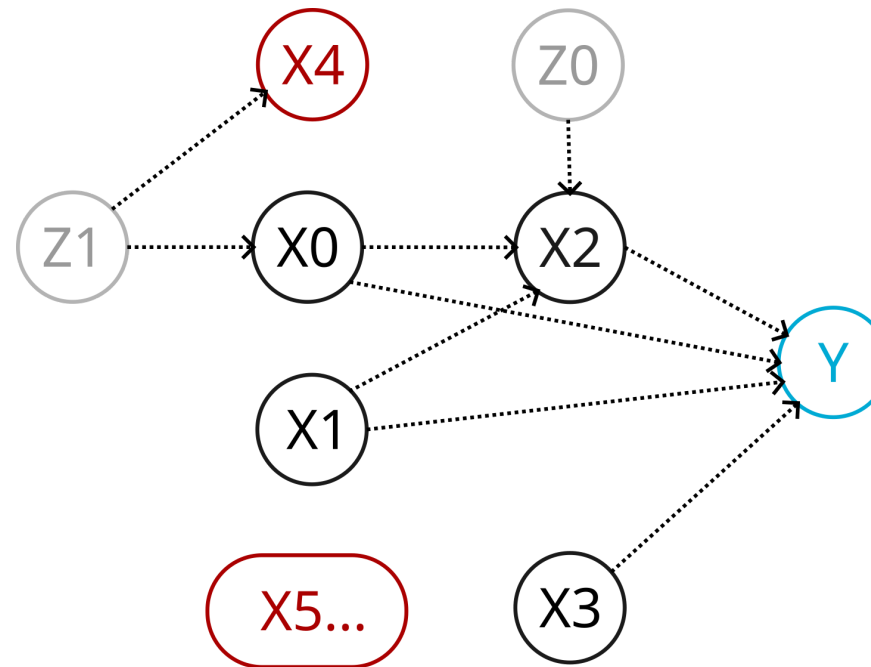
$X_2$  is also a complete mediator for  $Z_0$

- ...But in this case it is a good thing!
- $Z_0$  is not observed, but we can account for that at least indirectly



# Checking the Ground Truth

The process was engineered to contain several classical cases



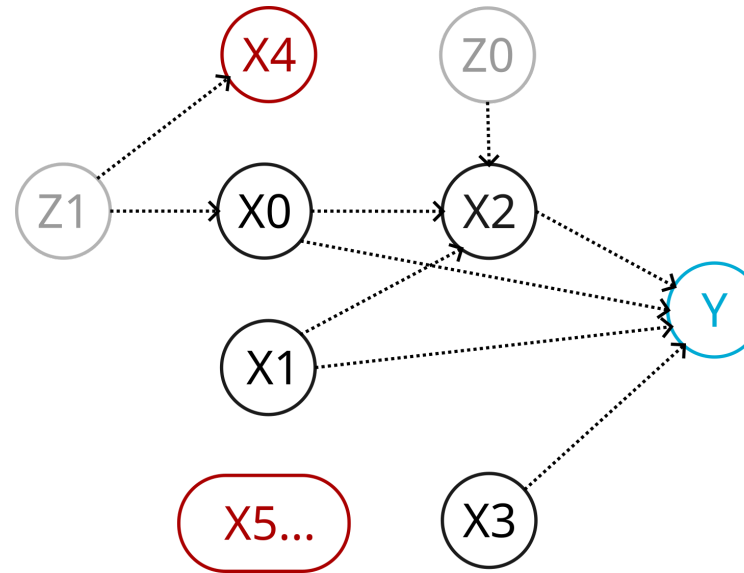
$Z_1$  is a **confounder** and causes a correlation between  $X_1$  and  $X_0$

- It is totally mediated by  $X_1$ , which is a good thing
- ...But it also causes a correlation between  $X_0$  and  $X_4$
- This might trick a model into considering  $X_4$  as important



# Checking the Ground Truth

Now let's check how accurate our importance estimate is:



```
In [65]: print(f'The accepted feature are {bfs.accepted}')  
print(f'...Which correspond to {[name_map[f] for f in bfs.accepted]}')
```

```
The accepted feature are ['u1', 'u12', 'u10', 'u13']  
...Which correspond to ['X1', 'X2', 'X3', 'X0']
```

If everything went as a planned, we should have found all relevant variables!



# A Few Final Remarks

## ML models are not just for prediction!

- They can be used for generation, anomaly detection, decision support
- ...And also as tools for a scientific analysis!

## Explainability is an important topic in AI

- It is one of the main approaches to make an AI model transparent
- This critical when AI systems need to interact with human users
- ...And for some domains it is also required by existing regulations

## Beware of correlated features

- Strongly correlated features (e.g mediated-mediator) may mislead algorithms
- Dealing with those is still a partially open problem!

